
StreamLip: Streaming Lip-to-Speech Synthesis via Joint Text-Audio Decoding with Flow Matching

Team Leader
2023001001

Team Member 1
2023002002

Team Member 2
2023003003

TL;DR

We propose **StreamLip**, the first streaming lip-to-speech (L2S) synthesis system that generates intelligible speech from silent video in real time. Our model jointly decodes text tokens and audio features within a unified autoregressive framework: a causal Transformer backbone with a look-ahead window predicts text tokens via autoregressive cross-entropy, while a lightweight flow matching head generates audio codec latents conditioned on both the visual features and the predicted text hidden states. This design addresses the fundamental viseme-to-phoneme ambiguity in L2S through an internal language model prior, without requiring an external lip-reading module at inference time. We plan to evaluate on LRS3 and demonstrate that our streaming approach, despite operating under strict causal constraints with only 200 ms of future context, can produce speech of reasonable intelligibility and naturalness.

1 Introduction

Lip-to-speech (L2S) synthesis aims to generate natural speech from silent talking face videos. This technology has potential applications in assistive communication for individuals with speech impairments, bandwidth-efficient video conferencing, and privacy-preserving voice reconstruction. Recent advances in diffusion and flow matching models have significantly improved the quality of L2S systems. LipVoicer [1] demonstrated that text guidance is essential for resolving the one-to-many viseme-to-phoneme mapping, reducing word error rate (WER) from 86% to 21% on LRS3. V2SFlow [2] introduced rectified flow matching to L2S and achieved naturalness surpassing ground truth. Most recently, SLD-L2S [3] established a new state of the art by operating in neural audio codec latent spaces.

However, all existing L2S methods share a critical limitation: **they operate offline**, requiring access to the entire video before generating any speech. This precludes real-time applications where low-latency speech output is necessary—such as live assistive communication devices that must maintain natural conversational rhythm (≤ 200 ms latency).

Meanwhile, adjacent fields have made progress on streaming generation. SoundReactor [4] proposed the first frame-level online video-to-audio generator using a causal autoregressive (AR) Transformer with a diffusion head, achieving 26 ms per-frame latency on game audio. Moshi [5] introduced the *Inner Monologue* mechanism for full-duplex spoken dialogue, where predicting time-aligned text tokens before audio tokens serves as a semantic scaffold that improves speech quality. UniVoice [6] unified AR-based ASR and flow-matching-based TTS within a single language model, demonstrating that discrete text prediction and continuous audio generation can coexist through careful loss balancing ($\lambda=0.005$) and dual attention mechanisms.

Our key observation is that these three ideas—streaming causal generation (SoundReactor), text-as-semantic-prior (Moshi), and AR+FM joint training (UniVoice)—can be synthesized into a unified framework for streaming L2S. We propose **StreamLip**, which makes three contributions:

1. **First streaming L2S system.** We formulate L2S as a causal, chunk-based generation problem with a 200 ms look-ahead window, enabling real-time speech synthesis from video.
2. **Joint text-audio decoding.** A single AR backbone simultaneously predicts text tokens (providing a language model prior to resolve viseme ambiguity) and conditions a flow matching head for audio generation, eliminating the need for an external lip-reading module.
3. **AR + flow matching co-training.** We adapt techniques from multi-task learning—stop-gradient isolation and scheduled corruption—to stably train the discrete text branch and continuous audio branch end-to-end.

2 Goal of Achievement

We aim to demonstrate a working streaming L2S system with the following measurable targets on the LRS3-TED benchmark [7]:

- **Intelligibility:** WER $\leq 40\%$ (measured via a pre-trained ASR model). For reference, LipVoicer [1] (offline, with external VSR) achieves 21%; V2SFlow [2] (offline, no text) achieves 28.5%.
- **Naturalness:** UTMOS ≥ 3.0 (ground truth is 3.52). V2SFlow achieves 3.78 in offline mode.
- **Latency:** End-to-end latency ≤ 250 ms (200 ms look-ahead + ≤ 50 ms compute). Moshi operates at ~ 200 ms for comparison.
- **Ablation evidence:** Demonstrate through controlled experiments that (a) the text branch improves WER over a text-free baseline, and (b) the look-ahead window improves quality over a purely causal baseline.

We consider WER $\leq 40\%$ a realistic target because: (1) the AR text head has access to full past context plus language model priors, which partially compensates for the limited visual window; (2) many WER “errors” on homophones (buy/by, their/there) produce acoustically identical speech.

3 Methods

3.1 Problem Formulation

Given a silent video sequence $V = \{V_1, \dots, V_n\}$ at 25 fps, we generate speech audio in a streaming fashion. At each time step i , the model can observe $V_{\leq i+\Delta}$ where $\Delta=5$ frames (200 ms look-ahead) and all previously generated audio. The output is a sequence of audio chunks, each corresponding to one video frame’s worth of speech (~ 40 ms).

Formally, we factorize the joint distribution as:

$$p(\mathbf{a}_{1:n}, \mathbf{w}_{1:n} \mid V_{1:n}) = \prod_{i=1}^n p(\mathbf{w}_i \mid V_{\leq i+\Delta}, \mathbf{w}_{<i}) \cdot p(\mathbf{a}_i \mid V_{\leq i+\Delta}, \mathbf{w}_{\leq i}, \mathbf{a}_{<i}) \quad (1)$$

where $\mathbf{w}_i$ is the text token at step i (predicted autoregressively) and $\mathbf{a}_i$ is the audio latent at step i (generated via flow matching conditioned on $\mathbf{w}_{\leq i}$ and $V_{\leq i+\Delta}$). Text is predicted *before* audio at each step, following Moshi’s Inner Monologue principle [5].

3.2 Architecture

The system consists of four components (Figure ??):

(a) Visual Encoder. We use a frozen AV-HuBERT (Large) [8] to extract visual features from cropped lip regions. AV-HuBERT outputs 1024-dim features at 25 Hz. To incorporate the look-ahead window, we apply chunk-level bidirectional self-attention: within each chunk of $1+\Delta$ frames, attention is bidirectional; across chunks, it is causal. A lightweight Conformer [9] adapter projects features to the backbone dimension.

(b) Causal AR Backbone. A decoder-only Transformer processes the interleaved sequence of visual features and token embeddings. At each time step i , the backbone receives the visual feature $\mathbf{v}_i$ (with look-ahead context) and the previous text token embedding $\mathbf{e}(\mathbf{w}_{i-1})$, producing a hidden state $\mathbf{h}_i$.

(c) Text Head (AR). A linear projection maps $\mathbf{h}_i$ to logits over the text vocabulary, trained with cross-entropy loss:

$$\mathcal{L}_{\text{AR}} = - \sum_{i=1}^n \log p(\mathbf{w}_i | \mathbf{h}_i) \quad (2)$$

During training, we use ground-truth text tokens from forced alignment. During inference, the text head autoregressively samples tokens. Critically, we pass the *continuous hidden state* $\mathbf{h}_i$ (not the discrete sampled token) to the FM head, preserving uncertainty information.

(d) Flow Matching Head (FM). A lightweight DiT [10] (4–6 layers) generates audio codec latents conditioned on $\mathbf{h}_i$ (which encodes both visual and textual information). Following SLD-L2S [3], we operate in the continuous latent space of a pre-trained audio codec (e.g., X-Codec or Mimi [5]).

The FM head models the vector field from noise $\mathbf{x}_0 \sim \mathcal{N}(0, I)$ to target audio latent $\mathbf{x}_1$:

$$\mathcal{L}_{\text{FM}} = \mathbb{E}_{t, \mathbf{x}_0, \mathbf{x}_1} \|v_\theta(\mathbf{x}_t, t | \mathbf{h}_i) - (\mathbf{x}_1 - \mathbf{x}_0)\|^2, \quad \mathbf{x}_t = (1-t)\mathbf{x}_0 + t\mathbf{x}_1 \quad (3)$$

where $t \sim \mathcal{U}(0, 1)$ and $\mathbf{h}_i$ is injected via adaLN-Zero conditioning [10]. At inference, we use an Euler solver with 10 steps and classifier-free guidance [11].

3.3 Training Strategy

Joint loss. The total loss combines both branches:

$$\mathcal{L} = \mathcal{L}_{\text{FM}}(\theta_{\text{fm}} | \text{stop_grad}(\mathbf{h}_i)) + \lambda \cdot \mathcal{L}_{\text{AR}}(\theta_{\text{backbone}}, \theta_{\text{text}}) \quad (4)$$

The FM loss gradient does not flow back through the backbone (stop-gradient), preventing the continuous audio gradients from destabilizing the discrete text predictions. Only $\mathcal{L}_{\text{AR}}$ updates the backbone. We set $\lambda=1.0$ for the AR branch since it is the sole driver of backbone learning.

Scheduled corruption. To prevent the FM head from over-relying on the text hidden states (which will be noisier at inference due to autoregressive error accumulation), we corrupt the text condition during training: with probability $p_{\text{mask}}=0.1$, we replace $\mathbf{h}_i$ with a zero vector; with probability $p_{\text{noise}}=0.05$, we add Gaussian noise to $\mathbf{h}_i$. This trains the FM head to gracefully degrade when text information is unreliable, falling back on visual features and acoustic continuity.

Training phases.

1. **Phase 1** (Backbone + Text Head): Train the visual encoder adapter and AR backbone on the lip-reading task only ($\mathcal{L}_{\text{AR}}$), warm-starting from a pre-trained LM if available. This produces a capable visual-conditioned language model.
2. **Phase 2** (Joint): Attach the FM head and train jointly with stop-gradient. The FM head learns to generate audio conditioned on the now-frozen-quality text hidden states.

3.4 Inference Pipeline

At each time step i :

1. Wait for frame $V_{i+\Delta}$ (200 ms look-ahead).
2. Encode $V_{[i-K, i+\Delta]}$ through the visual encoder with chunk-bidirectional attention.
3. The AR backbone produces $\mathbf{h}_i$ and samples text token $\mathbf{w}_i$.
4. The FM head generates audio latent $\mathbf{a}_i$ from noise, conditioned on $\mathbf{h}_i$, using 10-step Euler ODE.
5. Decode $\mathbf{a}_i$ through the pre-trained codec decoder to produce a waveform chunk.
6. Overlap-add with the previous chunk (50 ms overlap + linear crossfade) for smooth output.

3.5 Evaluation Plan

We evaluate on LRS3-TED [7] with the following metrics:

- **Intelligibility:** WER (via pre-trained ASR), SpeechBERTScore.
- **Quality:** UTMOS, PESQ.
- **Speaker similarity:** SECS (Speaker Embedding Cosine Similarity).
- **Lip sync:** LSE-C (Lip Sync Error - Confidence) via SyncNet.
- **Latency:** Per-frame waveform-level latency (ms).

Key ablations:

- **StreamLip (full) vs. w/o text head:** Measures the contribution of the language model prior.
- **StreamLip (full) vs. w/o look-ahead:** Measures the contribution of the 200 ms future context.
- **StreamLip (full) vs. w/ GT text:** Upper bound on text prior quality.
- **StreamLip (full) vs. Chunked V2SFlow:** Offline method constrained to same chunk size (fair baseline).

4 Timeline

Week	Milestone	Deliverable
1–2	Data preparation & environment setup	LRS3 preprocessed; codec trained/loaded
3–4	Phase 1: AR backbone + text head training	Streaming lip-reading model (WER metric)
5–6	Phase 2: Attach FM head, joint training	End-to-end streaming L2S (audio output)
7–8	Ablation experiments & hyperparameter tuning	Full evaluation table
9	Analysis, writing, poster preparation	Final report + poster

5 Workload Distribution

Member	Responsibilities
Member A (Leader)	Overall architecture design; AR backbone & text head implementation; loss balancing
Member B	Visual encoder integration (AV-HuBERT); data pipeline; look-ahead windowing
Member C	FM head implementation; audio codec integration; evaluation pipeline & ablations
All	Literature review; experimental analysis; report writing; poster

References

- [1] Yochai Yemini, Aviv Shamsian, Lior Bracha, Sharon Gannot, and Ethan Fetaya. Lipvoicer: Generating speech from silent videos guided by lip reading. In *Proc. ICLR*, 2024.
- [2] Jeongsoo Choi, Ji-Hoon Kim, Jinyu Li, Joon Son Chung, and Shujie Liu. V2sflow: Video-to-speech generation with speech decomposition and rectified flow. In *Proc. ICASSP*, 2025.
- [3] Ji-Hoon Kim, Jeongsoo Choi, and Joon Son Chung. Sld-l2s: Hierarchical subspace latent diffusion for high-fidelity lip to speech synthesis. *arXiv preprint arXiv:2602.11477*, 2026.
- [4] Koichi Saito, Julian Tanke, Christian Simon, Masato Ishii, Kazuki Shimada, Zachary Novack, Zhi Zhong, Akio Hayakawa, Takashi Shibuya, and Yuki Mitsufuji. Soundreactor: Frame-level online video-to-audio generation. *arXiv preprint arXiv:2510.02110*, 2025.
- [5] Alexandre Défossez, Laurent Mazaré, Manu Orsini, Amélie Royer, Patrick Pérez, Hervé Jégou, Edouard Grave, and Neil Zeghidour. Moshi: A speech-text foundation model for real-time dialogue. *arXiv preprint arXiv:2410.00037*, 2024.

- [6] Wenhui Gao et al. Univoice: Unifying autoregressive asr and flow-matching based tts. *arXiv preprint arXiv:2510.04593*, 2025.
- [7] Triantafyllos Afouras, Joon Son Chung, and Andrew Senior. Lrs3-ted: A large-scale dataset for visual speech recognition. 2018.
- [8] Bowen Shi, Wei-Ning Hsu, Kushal Lakhotia, and Abdelrahman Mohamed. Learning audio-visual speech representation by masked multimodal cluster prediction. In *Proc. ICLR*, 2022.
- [9] Ankit Gulati et al. Conformer: Convolution-augmented transformer for speech recognition. In *Proc. Interspeech*, 2020.
- [10] William Peebles and Saining Xie. Scalable diffusion models with transformers. In *Proc. CVPR*, 2023.
- [11] Jonathan Ho and Tim Salimans. Classifier-free diffusion guidance. In *NeurIPS 2021 Workshop*, 2021.